From incident reports to decisions people can inspect

A bilingual workspace for evidence, draft communication and human review.



Keep facts, uncertainty and responsibility visible.
No automatic operational execution.

APPLICATION CODE VERSION

27f5ec9

Application code unchanged since 27f5ec9; later commits update documentation and submission materials.

Recorded evaluations and video keep their original checkpoints.

Benjamin Jin, class 9(4) · GitHub: zha-chicken

Project owner · designed the guided-workflow UI/UX and interaction logic

A risk signal is not a confirmed flight outcome

FOR EXAMPLE

A weather warning
≠ a flight delay

An operational suggestion
≠ authority to act

Teams must interpret incomplete or conflicting reports without turning uncertainty into a promise.



Fragmented evidence

Different report times, source reliability and historical comparisons.



Different audiences

Operations, airport coordination and passengers need different drafts.



Human responsibility

Safety, release, dispatch and publication decisions stay with authorized people.

“I value UI/UX as much as analysis logic.”

I’ve been suffering from bad UI/UX designs and bad interaction logics for a long time, so I know the importance of them. For example, I believe system leading the user to their wanted function is way better than the user finding the next function that they would want in a page crowded with functions

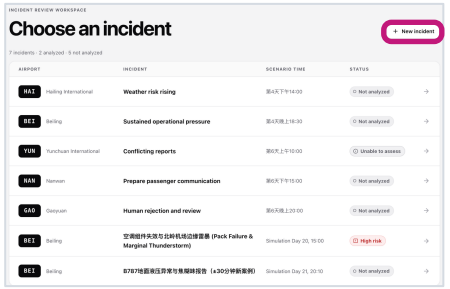
MY DES/GN

GUIDED WORKFLOW: The system leads each step

Before: every tool shared one tabbed incident page. **Now:** a home page, then one clear next step at a time.

01 · RECORD

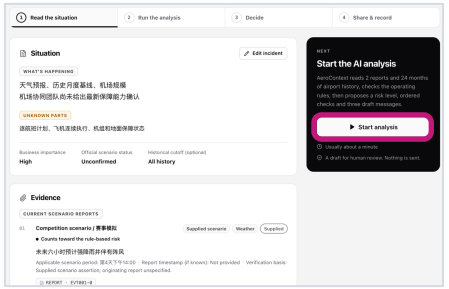
Home



Choose an incident from the list, or create a new one.

01 · RECORD

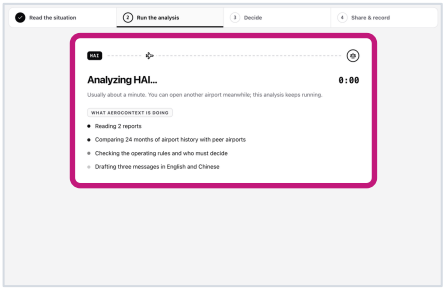
One next step



Read the situation and evidence; the one primary action is Start analysis.

02 · ANALYZE

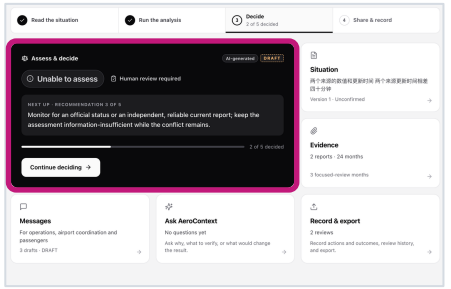
Focused analysis



While it runs, the main area shows only the analysis.

03 · DECIDE

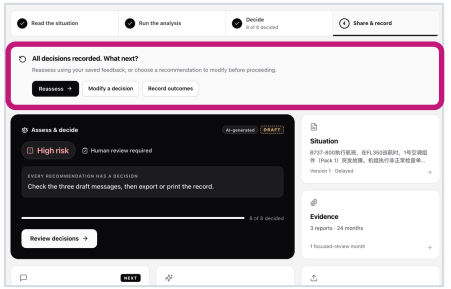
Six large blocks



The main area holds six large blocks, led by Assess & decide.

04 · REVISIT

What next?



After the last decision, a banner offers Reassess, Modify a decision or Record outcomes.

01 Record

Incident, reports, privacy review and source/time metadata

02 Analyze

Scoped history, deterministic policy, model explanation

03 Decide

Accept, modify, reject or leave Pending; retain full reasons. Accepted does not mean executed; Pending does not mean approved.

04 Revisit

Record outcomes later; reassess or edit; preserve revisions. An unrecorded outcome remains unknown.

One record from evidence to later outcomes

English and Chinese · clickable citations · versioned history · JSON export and print

Screens: real captures of the local prototype with fictional data (8ad14e3 build; UI unchanged in 27f5ec9). Ease of use and time savings were not measured.

SIMULATION PROTOTYPE

Calculate constraints first. Generate prose second.

The model receives scoped evidence and policy results; validation cannot prove every sentence true.

1 · CALCULATE

Deterministic context

- Current reports within ±30 minutes
- Historical mean: up to 12 prior months
- Source independence and stop conditions
- Known evidence / metric identifiers

2 · GENERATE

Model

- DeepSeek Flash / high / Responses
- Bilingual explanations and three drafts

3 · CHECK

Bounded validation

1× repair

- JSON, evidence and constraint checks
- One repair; visible failure or advisories

RECORDED REAL-MODEL EXAMPLE

SIMULATION PROTOTYPE

Gusts + an expanding support shortage

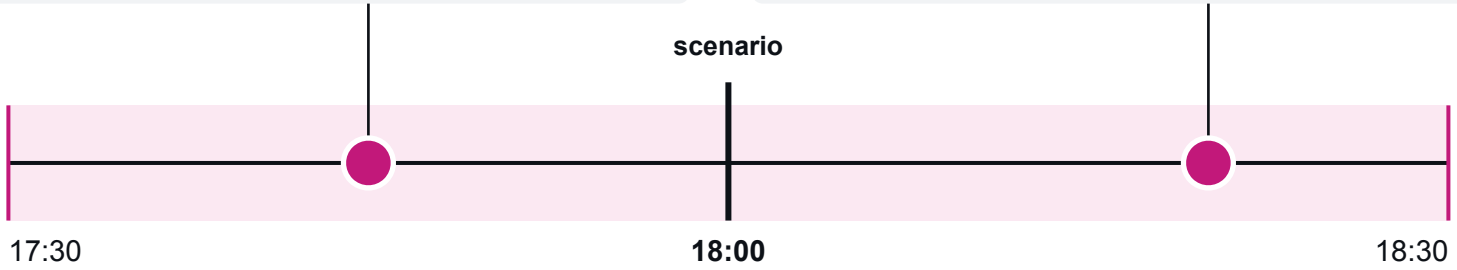
UNSEEN02-en • fictional QING • scenario Day 42, 18:00 • unseen when tested, now disclosed

17:45 Gust report

Verified in the fictional log.
Separate weather origin.
15 minutes before scenario.

18:20 Support report

Shortage spreads to a second stand cluster.
20 minutes after scenario.



±30-minute current-report window

RESULT

High + human review

Two independent signal types plus expanding impact.

No official flight delay.

THE REPLY PROPOSES

referral · verification · monitoring · drafts

RECORDED RESULT

one call	100.62 s	CNY 0.176188 estimated	qualitative notes retained

Source: evidence/unseen/cases/UNSEEN02-en — definition.json, outputs.md and evaluation.json

Usable output and good reasoning are different

One bounded run per set; no claim of statistical superiority or hallucination-free behavior.

Measure	Supplied original	Supplied rerun	Unseen
Accepted model steps	32 / 32	32 / 32	21 / 22
Usable final replies	28 / 28	28 / 28	19 / 20
Quality failures*	4 / 30	3 / 30	3 / 20 + 1 blocked
Steps needing repair	12 / 32	18 / 32	7 / 22
Median latency	67.741 s	81.907 s	69.273 s
Estimated cost, CNY	4.894533	5.553762	3.420530

UNSEEN · USABLE FINAL REPLIES

19 / 20

Accepted by the app; not proof the content is free from unsupported claims.

UNSEEN · QUALITY REVIEW

3 / 20 + 1 blocked

Recorded project review, not an independent human panel.

*Quality denominator: final language variants; supplied sets include two privacy checks without model calls.
Model-step totals include baseline assessments. Cost uses app accounting, not provider invoices.
Original supplied tests: 7bde4c8. Rerun and unseen tests: code later committed as 8ad14e3. No model reevaluation for 27f5ec9.

Sources: D03; evidence/supplied-original, supplied-current and unseen summary.json

Preserved decisions—even when reassessment fails

Five synthetic records • all four dispositions • seven revisions • two later simulated outcomes

Initial assessment saved
real model call; high risk

Five synthetic reviews
accepted · modified · rejected · pending

Reassessment failed
an uncited fact

Its one repair failed
invented a source identifier

Prior assessment kept
all reviews intact

What was demonstrated

- ✓ Original recommendations and reasons retained.
- ✓ Later outcomes added without losing history.
- ✓ Previous valid assessment preserved on error.

What was not demonstrated

- ✗ The reassessment failed after one repair.
- ✗ The repair invented a source identifier.
- ✗ These are not real participant observations.

Sources: D05-synthetic-feedback.md/.csv/.json; evidence/feedback/reassessment-review.md. Synthetic demonstration, not a real-user study.

Traceability is demonstrated.
Time savings are unmeasured.

The purpose is to make evidence and reviewer judgment easier to inspect.

Available evidence

- Recorded model replies, quality reviews and costs.
- 351 automated tests; TypeScript + Webpack passed.

The latest default-build attempt hit an environment worker-port restriction.

Accepted limitations

- Unsupported commitments.
- Historical interpretation errors.
- Validator false positives.
- Input privacy is heuristic.

Intentionally skipped

- Manual timing study.
- Basic-prompt comparison.
- High-versus-max experiment.
- Standalone B787 package.

Time savings and superiority over a basic prompt were not measured.

A reviewable package before a supervised release

Eight required materials and the approved video are linked from the repository.

Review now

- Editable presentation and technical report
- Source, rules, tests and frozen evidence
- Synthetic feedback and risk card
- Approved video and source hashes

Before submission

OPEN GATES

- [ ] Owner makes repository public
- [ ] Verify links and submit repository URL
- [ ] Rehearse; retain portal confirmation
- **Fresh judge install and Docker still unverified**

No hosted deployment or public access is claimed.

PROJECT OWNER

Benjamin Jin, class 9(4)
Designed the guided-workflow UI/UX
GitHub: zha-chicken
github.com/zha-chicken/runway-incident-workspace

AI USE

OpenAI Codex and Anthropic Claude assisted development and documentation; Claude created the owner-approved video. Recorded app replies use DeepSeek Flash.

DATA

Organizer-supplied Track 3 Problem Statement (problem setters: Accenture) and Track 3 Data.xlsx. Airports and incidents are fictional; D05 records are synthetic.

Full credits and AI-use statement: ACKNOWLEDGEMENTS.md. Sources: submission/README.md; D07; D08; publication verification.